


12-HOUR ENGINEERING HACKATHON

Teach the phone to listen.

Build the AI that recovers a dropped customer journey — over a real phone call — and knows exactly when to hand a frustrated caller back to a human.



12
HOURS

01
SOLO BUILD

1
VERTICAL

On-site
JAIPUR OFFICE

What you're walking into

CIMET runs comparison journeys across seven verticals. Customers arrive from our own econnex site, from campaigns, and from affiliates on their own platforms. Each journey is a multi-step form that collects what's needed to match a customer to a product.

Energy

Broadband

Mobile & Internet

Personal Loans

Credit Cards

Solar

Private Health

HACKATHON SCOPE

This build is scoped to a single vertical — **Energy**. Everything you design and demo should go deep on the Energy journey end-to-end; the other six verticals are context, not build targets, for this session.

Plenty of customers don't finish in one sitting. The moment someone enters a phone number and email, then abandons a later step, they become a dropped-off lead — and recovering that lead is real revenue.

How recovery works today

It works. It also depends entirely on a human reading scripts and typing answers, one call at a time. Agents operate in two modes — and that duality is your solution space.

MODE A

Agent-Assisted

A human guides the customer, supported by the console.

MODE B

Agent-Driven

The agent runs the journey on the customer's behalf.

01

Cron picks the lead

Scheduled jobs collect dropped-off leads and pass them to ViciDial, our dialler.

02

Agent calls back

The agents team rings the customer from the agent console.

03

Reads the script

Every section and field has a predefined script the agent reads to seek data.

04

Fills the journey

The real journey runs in an iframe; the agent completes it for the customer.

Bring AI to the call — on two fronts

Voice is the primary target this round. The console upgrade is a good-to-have extension if you have time left over — but a strong submission is judged first on how well it handles the call itself.

01

PRIMARY

Let a voice agent run the call end-to-end

For suitable cases, can an AI voice agent place the call, hold the conversation, seek the data, complete the journey on its own — and escalate to a human when it should? This is Agent-Driven mode, automated. This is where we want most of your 12 hours to go.

02

GOOD TO HAVE

Modernise the console so agents fly

The console today is a manual tool. In an AI-native version, how much of the agent's cognitive load can the system carry — surfacing the right script at the right moment, capturing answers as the customer speaks, cutting the typing and the searching? The measure of success is agent efficiency and quality: more journeys completed, faster, with fewer errors.

Where to start

Begin with dropout recovery. It's the most tractable scenario: the customer already started a known journey, the path is bounded, the stakes are lower than a retention or cancellation call, and the data to collect is well-defined. Prove it there — for Energy — then, time permitting, reach for the harder cases.

Know when to hand back the call

An autonomous voice agent is only trustworthy if it knows its own limits. A core capability of any voice solution is recognising — in real time — when to transfer to a live human, with a warm handoff that passes the context and everything collected so far, so the customer never repeats themselves.

RECOVERY CALL • TRANSCRIPT

● REC

AI AGENT Hi Priya — you started comparing energy plans with us yesterday. Mind if I pick up where you left off? This call is recorded.

CUSTOMER Look, I've already told three of you my details. This is the second call today.

CUSTOMER Honestly I'm done with this — just stop calling me.

↳ FRUSTRATION DETECTED — WARM HANDOFF

HUMAN Hi Priya, this is Aarav — I can see everything so far, no need to repeat anything. Let me sort this out for you.

Escalate on signals such as

ANGER

Raised voice, profanity, repeated complaints, or sentiment turning negative.

CONFUSION

The agent mishears or fails to capture the same field two or three times.

OFF-SCRIPT

Questions the scripts don't cover, or requests for advice the agent can't give.

SENSITIVE

Anything touching payment, disputes, or a vulnerable customer.

ASKS

The customer explicitly requests a person.

LOW CONF

The agent is simply unsure what was said or what to do next.

How you detect and route is up to you. That it does so reliably, and hands over gracefully, is part of what we judge.

Scripts are yours to design

Today, human agents already work from a script: every section and field has a predefined line they read to seek exactly the right data.

ASSUME, THEN DESIGN

Assume the same structure exists on the human-agent side — one script per section, one per field, refined to collect exactly what's needed. Your job is to draft your own dummy scripts that capture that same intent, based on the call recording we give you and the field list for the Energy vertical.

Whether you assist a human or drive an autonomous call, treat the scripts you design as the source of truth for what the conversation collects. Reasonable, well-structured scripts are part of the build — not a shortcut you're handed. Make the seeking and capturing intelligent and automatic on top of whatever you draft.

A strong solution shows

Go deep on one path, not shallow on many. One vertical, one scenario, done convincingly beats several half-built ones.

- A real journey — Energy, end-to-end — completed through your system, submitting valid data the way the console does today.
- Your own drafted scripts, reused consistently to drive what's asked and captured.
- For voice (primary): a conversation that survives real human behaviour — interruptions, mishears, "I'm busy," "not interested" — and hands off to a human, with context, the moment things turn.
- For the console (good-to-have): a measurable cut in the manual work an agent does per call.
- Clear evidence of an efficiency or quality gain over the manual workflow.

Constraints & guardrails

We operate in regulated Australian verticals. For a prototype these can be simulated — but the behaviour must be designed in.

▶ Test data only No real customer PII, no real numbers. Use the synthetic leads and test numbers we provide.	▶ Consent first Any call discloses recording before collecting anything (AU two-party consent).
▶ No card data by voice Never capture payment details through an AI voice channel — hand off instead. Show the boundary.	▶ No advice The system collects information; it does not give product or financial advice.
▶ Do-Not-Call aware Assume a DNC check gates dialling (ACMA register). You may stub it — but show where it sits.	▶ Respect "no" If a customer declines, thank them, log it, end. No pressure loops.

How we'll score it

CRITERION	WEIGHT	WHAT WE LOOK FOR
Working outcome	30%	A journey actually completed, valid data submitted — live, in front of judges.
Efficiency / quality gain	25%	Real reduction in manual effort or improvement in completion vs. today.
Robustness & handoff	20%	Handles messy human conversation, and hands to a human — gracefully, with context.
Scripts & integration	15%	Built on your own designed scripts; correct payload shape; sensible escalation.
Guardrails & judgment	10%	Consent, payment boundary, DNC awareness, respecting declines.

The demo is a live run — a real call or assisted journey, narrated. Slides support it; they don't replace it.

What we hand you

You choose your own stack, models, and tools beyond what we provide.

- A test recording of a real recovery call — hear exactly how our agents run it today: pacing, tone, objections handled, data sought. Build your solution based on what you hear.
- A synthetic lead dataset for Energy — ID, last-completed step, test contact.
- A journey-completion sandbox or mock, with the expected payload shape.
- Stable, dedicated internet connectivity on-site.
- On-site mentors from the platform and agents teams.

What you bring

This is a bring-your-own-tools build. CIMET provides the problem, the data, connectivity, and the mentors — everything you build with is on you.

- Your own laptop, charger, and any accessories you need to work through the full 12 hours.
- Your own access to whichever AI tools you build with — Claude, Cursor, ChatGPT, or otherwise.
- Your own subscriptions or licenses for any tool, model, or service your solution depends on.
- Anything else your chosen stack requires — CIMET does not provision or reimburse tooling for this session.

COME READY TO BUILD

Arrive with your environment already set up and your logins already working — the clock starts at hour zero, and setup time comes out of your 12 hours.

HOW TO ENTER

One hard problem. 12 hours. Solo.

We're inviting strong developers in Jaipur. Bring your laptop and your best engineering instincts — we'll bring the problem, the data, and the mentors.

[APPLY VIA LINKEDIN →](#)

Apply by [date] · Limited spots

12 hrs

FORMAT · ON-SITE

Solo

TEAM SIZE — 1

Jaipur

CIMET OFFICE

Repo + demo

SUBMISSION

The ask, one more time:

Read the room. Know when to step aside.

Then recover the journey.