


QA AUTOMATION • ENGINEERING BRIEF

Score the sale before it ships.

Every sale in CRM gets checked against its retailer's checklist automatically, before it can submit. All green, it goes through untouched. Any red, it's held — with the exact failing line, transcript, and timestamp sitting right next to it.



12	01	3	On-site
HOURS	SOLO BUILD	RETAILERS	JAIPUR OFFICE

A manual process that can't scale

An auditor pulls a recording, listens end to end, fills an Excel form, and the sale waits. We do thousands of sales a month across roughly 30 retailers. This doesn't scale, and it won't survive a serious retailer audit.

THE ASK

Every sale in CRM is scored automatically against that retailer's checklist before it can be submitted. All green, the sale goes through untouched. Anything red, the sale is held and the agent or TL sees exactly which criteria failed, with the transcript line and audio timestamp sitting right next to it. They fix it or challenge it, then it goes.

Three of our current forms are already attached — energy, broadband, plus the retailer-side forms from Retailer 1, Retailer 2 and Retailer 3. These are already the spec. We just need to move them out of Excel and into CRM.

FORMS ALREADY IN HAND

Energy

Broadband

Retailer 1

Retailer 2

Retailer 3

Recording and transcript ingestion. Zero manual handling.

The dialler pushes the recording to CRM by API, keyed on Lead ID, within minutes of hangup. CRM sends it to transcription with speaker separation and word-level timestamps. Transcript is stored against the lead. Nobody downloads or uploads anything, ever.

BUILD THIS FIRST

Everything else in this brief — the check library, the scoring engine, the dashboards — is dead without this pipeline. Recording and transcript ingestion, keyed on Lead ID, with zero manual handling, comes before anything else.

Once a transcript exists against a lead, scoring becomes a comparison problem: transcript against script, transcript against CRM fields, transcript against the retailer's own plan and rate card. That comparison is what the rest of this brief specifies.

This part matters most

TYPE	COMPARES	EXAMPLE	BLOCKS SALE?
A • Verbatim / script	Transcript vs. approved script	DMO/VDO read verbatim, recording disclaimer, T&Cs, EIC	Yes if critical
B • Factual match	Transcript vs. CRM fields vs. retailer plan and rate card	Rate quoted = rate on plan; address, DOB, email, NMI/MIRN, fuel type, concession, life support, move-in date, gift card value	Yes
C • Behaviour	Transcript only	Rapport, interruptions, dead air, objection handling	No

Lead 3613790 — illustrative

Retailer 1, Agent A, 30 minute call. CRM shows:

LEAD 3613790 • RETAILER 1 • 30 MIN CALL

- ✓ **PASS** Recording disclaimer (critical) — verbatim match at 00:12
- ✓ **PASS** Account holder confirmed (critical) at 02:41
- ✓ **PASS** DMO read verbatim (critical) at 09:55
- ✗ **FAIL** Rates and charges (critical) — agent said "peak is 28.6 cents" at 14:02. Plan on the lead is 31.9c/kWh peak. Mismatch.
- ✗ **FAIL** Email captured (critical) — agent read back "j.smith@gmail.com" at 22:10. CRM field says "j.smith@gmial.com." Mismatch.
- **NOTE** Dead air (non-critical): 47 seconds at 18:30. Coaching note.

Sale is held. The TL opens the lead, clicks the timestamp, hears twenty seconds, corrects the email and either re-confirms the rate with the customer or the sale is cancelled. Nobody listened to thirty minutes to find that.

When a sale actually goes through

- All criticals pass — sale auto-submits, no human touch.
 - Any critical fail — held, routed to the TL queue.
 - Low confidence on any check — routed to QA rather than auto-passed.
 - 5% of clean calls are sampled to a human anyway — we're measuring the model, not just trusting it.
-

Agent-wise, daily / weekly / monthly

- Score with and without fatal factors (we already calculate both).
- Critical fail rate, and which specific check is failing.
- First-pass yield — percentage of sales that go green with no rework.
- Repeat offences — the same critical check failing three or more times in a rolling seven days flags the TL and triggers a warning under the performance policy.

Rolled up by retailer, campaign (owned site, affiliate, paid search, inbound), site, and TL. Plus auditor agreement rate, so we can see whether the model is calibrated and whether our own auditors are consistent with each other.

Constraints & guardrails

We operate in regulated Australian verticals. The scoring system inherits the same obligations as the calls it's grading.

▶ **Test data only**

No real customer PII in any dev or demo environment. Use the synthetic leads and sanitised transcripts we provide until the gate is live on real traffic.

▶ **Consent is a check, not an assumption**

The system verifies the recording disclaimer was read on every call; it never assumes consent just because a recording exists.

▶ **No card data surfaced**

If a card number is spoken on a call, the transcript view redacts it before anyone sees it. The score flags the violation; the number itself never needs to be shown.

▶ **No advice, no auto-correction**

The system tells you what failed and exactly where. It doesn't rewrite the sale, override the agent, or contact the customer on its own.

▶ **Scored against the rules that were live**

Every check resolves to the check-library version in effect on the call date, not today's version.

▶ **Respect the override**

When a human auditor overturns an automated score, that decision is logged, not silently dropped.

How we'll score it

CRITERION	WEIGHT	WHAT WE LOOK FOR
Scoring accuracy	30%	Agreement with human auditors on a sampled set of calls; critical checks essentially never false-pass.
Coverage	25%	All three check types — verbatim, factual, behaviour — running end-to-end against at least one retailer's full checklist.
Traceability	20%	Every score resolves to a transcript line, an audio timestamp, and the exact check version that was live on the call date.
Gate & escalation logic	15%	Any critical fail is held, low-confidence scores route to a human, nothing uncertain auto-passes.
Guardrails & judgment	10%	Handles messy real calls — crosstalk, mishears, silence — without manufacturing false criticals.

The proof is a real lead scored live — held or passed, with the exact reasoning shown. Slides support it; they don't replace it.

What we hand you

You choose your own stack, models, and tools beyond what we provide.

- A test recording of a real sales call — hear exactly how agents talk through each section, so your comparison logic is built against real speech, not assumptions.
- A synthetic lead dataset — ID, retailer, last-completed step, test contact, and a sanitised transcript to develop against.
- A check-library export for one retailer — criteria, weights, and critical flags already defined, ready to load into your scoring engine.
- A scoring sandbox with the expected payload shape, so you're not guessing the schema.
- Stable, dedicated internet connectivity on-site.
- On-site mentors from the platform and QA teams.

What you bring

This is a bring-your-own-tools build. CIMET provides the problem, the data, connectivity, and the mentors — everything you build with is on you.

- Your own laptop, charger, and any accessories you need to work through the full 12 hours.
- Your own access to whichever AI tools you build with — Claude, Cursor, ChatGPT, or otherwise.
- Your own subscriptions or licenses for any tool, model, or service your solution depends on.
- Anything else your chosen stack requires — CIMET does not provision or reimburse tooling for this session.

COME READY TO BUILD

Arrive with your environment already set up and your logins already working — the clock starts at hour zero, and setup time comes out of your 12 hours.

HOW TO ENTER

One hard problem. 12 hours. Solo.

We're inviting strong developers in Jaipur. Bring your laptop and your best engineering instincts — we'll bring the problem, the data, and the mentors.

[APPLY VIA LINKEDIN →](#)

Apply by [date] · Limited spots

12 hrs

FORMAT · ON-SITE

Solo

TEAM SIZE — 1

Jaipur

CIMET OFFICE

Repo + demo

SUBMISSION

The ask, one more time:

No sale ships unscored.

Then explain exactly why, in seconds — not minutes.